MACHINE LEARNING COURSE PROJECT PROPOSAL • FALL 2026

Risk-Aware Machine Learning Under Distribution Shift

Epistemic Uncertainty Quantification & Selective Classification in Financial Tabular Streams

PROJECT TEAM — 3 MEMBERS

Jay Lakkad	Hiren Modhvadia	Dev Pandya
M26CSE005	M26CSA012	M26CSE003

Quadrant 1 • Problem, Motivation and Impact

FOUNDATIONS & SCOPE

THE CORE PROBLEM

Standard supervised ML relies on the **i.i.d. assumption** ($P_{train}(X,Y) = P_{test}(X,Y)$). In high-stakes financial domains (fraud detection, credit scoring), this assumption fails due to evolving attacker tactics, macroeconomic swings, and temporal shifts. High-capacity models make **catastrophic, overconfident errors** on shifted data without self-awareness of their unreliability.

LIMITATIONS OF CURRENT PRACTICE

- **Poor Calibration Under Shift:** Raw softmax or tree leaf probabilities are uncalibrated out-of-distribution (OOD), yielding high confidence on incorrect guesses.
- **Noise vs. Novelty Conflation:** Conventional models fail to separate aleatoric noise (feature overlap) from epistemic uncertainty (novel distribution spaces).
- **OOD ≠ Decision Risk:** Distance-based OOD detectors flag unfamiliar points regardless of whether the prediction itself is safely separable.

BENEFICIARIES

Fraud Analysts & Risk Teams: Workload reduced by auto-clearing safe transactions and routing ambiguous ones to triage.
Consumers: Minimized false card declines and zero-day fraud losses.

SCIENTIFIC & PRACTICAL IMPACT

Formalizes selective classification for tabular data, enabling a **risk-aware human-in-the-loop** paradigm that slashes shift-induced fraud losses while automating >85% of volume.

Quadrant 2 • Dataset & Controlled Shift Suite

DATA BENCHMARK

Attribute	Benchmark Details (ULB Credit Card Fraud Benchmark)
Data Type & Modality	Tabular financial transactions (continuous numerical features)
Dataset Size	**284,807 transactions** recorded over 48 hours; 30 input features
Classes & Imbalance	Binary: **492 frauds (0.172%)** vs 284,315 legitimate (99.828%)
Features	`Time`, `Amount`, and 28 anonymized PCA components (`V1` – `V28`)
Working Link	kaggle.com/datasets/mlg-ulb/creditcardfraud (Dal Pozzolo et al., 2015)

REPRESENTATIVE DATA SAMPLE SNIPPET

Time(s)	V1	V2	V3	Amount(€)	Class	Protocol Note
0.0	-1.3598	-0.0728	2.5363	149.62	0	Normal Legitimate (ID)
541.0	-2.3122	1.9519	-1.6099	0.00	1	Actual Fraud Record
472.0	-3.0435	3.1573	-3.1674	529.00	1	High-Value Attack (Shift)

CONTROLLED SHIFT INJECTION PROTOCOL

Split: Chronological 70% Train (199k), 15% Calib/Val (42k), 15% ID Test (42k).
Shifts: (1) *Covariate Shift:* Gaussian jitter on key features V1–V4; (2) *Subpopulation Shift:* High-amount tail strata (>90th pct); (3) *Prevalence Shift:* Skew fraud rate 0.17% → 2.5%; (4) *Temporal Drift:* Final 8-hr slice. *Ethics:* Anonymized PCA, zero PII.

Quadrant 3 • Proposed Work & Key Contributions

METHODOLOGY

1. Base Classifiers (LogReg, RF, LightGBM, Tab-MLP) → **2. Uncertainty Engine** (Isotonic, Conformal, Ensemble Var) → **3. Risk Scoring** (Risk = f(Confidence, Disagreement)) → **4. Gated Decision** (Accept (≤τ) vs Flag for Review (>τ))

CORE ALGORITHMS & BASELINES

- **Predictors:** LightGBM, Random Forest, Logistic Regression, Tabular Deep Ensemble.
- **Uncertainty Estimators:** (a) *Calibration:* Isotonic Regression & Platt Scaling; (b) *Ensemble Disagreement:* Inter-tree variance $U_{ens} = (1/M)\Sigma(\hat{p}_m - p)^2$; (c) *Conformal Prediction:* Non-conformity sets with coverage guarantee $1 - \alpha$.
- **Baselines:** (1) Vanilla Softmax with static 0.5 threshold; (2) Random deferral; (3) Isolation Forest OOD filter.

KEY PROPOSED CONTRIBUTIONS (SPECIFIC & DISTINCT)

- **C1. Tabular Shift Benchmark Suite:** Standardized multi-axis synthetic & temporal shift generation harness for imbalanced tabular benchmarks.
- **C2. Comparative Risk Estimator Bake-off:** First direct empirical comparison of calibration vs. ensemble variance vs. conformal bounds under tabular shift.
- **C3. Cost-Sensitive Selective Policy:** Risk-Coverage optimization routing ambiguous decisions to human analysts while preserving ≥85% automated clearance.
- **C4. Shift Detection vs. Decision Risk Disentanglement:** Clarifying that OOD feature status does not imply misclassification risk, and vice versa.

Quadrant 4 • Comprehensive Evaluation Plan

VALIDATION & METRICS

EVALUATION PROTOCOL & SHIFT TESTBEDS

Trained exclusively on chronologically early In-Distribution (ID) transactions (N = 199,365). Evaluated across ID holdout and 4 distinct shifted test partitions under varying perturbation intensities $\varepsilon \in [0.1, 2.0]$.

AUPRC & AUROC
Discriminative power under 0.17% fraud prevalence

Expected Calib. Error
$ECE = \Sigma (|B_k|/N)|acc(B_k) - conf(B_k)|$

AURC (Risk-Coverage)
Area Under Risk-Coverage for c ∈ [0.6, 1.0]

Cost-Weighted Loss
L = $100•FN + $5•FP + $2•Defer

ABLATION & ROBUSTNESS MATRIX

- **Shift Intensity:** Monitor degradation curves across severity scales $\sigma \in \{0.2, 0.5, 1.0, 2.0\}$.
- **Rejection Budget Trade-off:** Sweep deferral fraction from 0% to 25% to map the Pareto frontier.
- **Uncertainty Decomposition:** Compare calibration alone vs ensemble variance vs hybrid risk gating.

✓ Concrete Success Criteria

1. **≥30% Error Reduction:** Selective abstention cuts shift-induced loss by ≥30% over vanilla baselines.
2. **High Throughput:** Retains **≥85% automated decision coverage** under moderate shift.
3. **Calibrated Reliability:** Post-calibration ECE < 0.05 across shifted test sets (50% improvement).